

PYTHON MISC

```
convert row vector into column vector  
Y = Y.reshape(-1,1)
```

sigmoid

```
from scipy.stats import logistic  
logistic.cdf(x)
```

LINEAR CLASSIFIERS	STOCHASTIC GRADIENT DESCENT	KERNEL PERCEPTRON	MISC
boundary line: $\theta \cdot x + \theta_0 = 0$	update: sample i at random	$\theta = 0$ for $i = 1 \dots n$ <div> if $y^{(i)} \cdot \theta \cdot \phi\left(x^{(i)}\right) \leq 0$ <div> $\theta \Leftarrow \theta + y^{(i)} \phi\left(x^{(i)}\right)$ </div> </div>	$\vec{a} \cdot \vec{b} = \sum a_i b_i$ $a \cdot a = \ a\ ^2 \Rightarrow$ $\theta_z \cdot Ax = \theta_z^T Ax = \left(A^T \theta_z\right) \cdot x$
PERCEPTRON	$\theta \Leftarrow \theta - \eta_T \nabla_{\theta} \left[Loss_h\left(y^{(i)} \theta \cdot x^{(i)}\right) + \frac{\lambda}{2} \ \theta\ ^2 \right]$	$\theta \cdot \phi\left(x^{(i)}\right) = \sum_j^n \alpha_j y^{(j)} \phi\left(x^{(j)}\right) \phi\left(x^{(i)}\right)$ where α_j is nb of mistakes made on j th example	distance from point to line $d = \frac{ \theta \cdot x + \theta_0 }{\ \theta\ }$
$y^{(i)} \in \{-1, 1\}$	$= \theta \left[\begin{cases} 0 & \text{if } Loss_h = 0 \\ -y^{(i)} \cdot x^{(i)} & \text{if } Loss_h > 0 \end{cases} + \lambda \theta \right]$	we can then rewrite the perceptron algorithm without having to transform $x \Rightarrow \phi(x)$ anymore	derivative of norm squared $\left(\ h(x)\ ^2\right)' = 2h(x)h'(x)$
algorithm $\theta = \vec{0}$ for t = 1...T for i = 1...n if $y^{(i)}\left(\theta \cdot x^{(i)} + \theta_0\right) \leq 0$ then $\theta = \theta + y^{(i)} x^{(i)}$ $\theta_0 = \theta_0 + y^{(i)}$	LINEAR REGRESSION empirical risk $R_n(\theta) = \frac{1}{n} \sum \frac{\left(y^{(t)} - \theta x^{(t)}\right)^2}{2}$ small deviations less penalised than large ones $\nabla_{\theta} R_n(\theta) = -\left(y^{(i)} - \theta x^{(t)}\right) \cdot x^{(t)}$ gradient descent $\theta = 0$ for: random $t = \{1 \dots n\}$ # stochastic, could also loop $\theta = \theta + \eta_k \left(y^{(t)} - \theta \cdot x^{(t)}\right) \cdot x^{(t)}$ with $\eta_k = \frac{1}{1+k}$	$\theta = 0$ for $i = 1 \dots n$ <div> if $y^{(i)} \sum_j^n \alpha_j y^{(j)} K\left(x^{(j)}, x^{(i)}\right) \leq 0$ <div> $\alpha_i = \alpha_i + 1$ </div> </div>	derivative $[y(\theta \cdot x)]' = yx$ because we derive for each θ_i
REGULARIZATION		kernel composition $\cdot f: \mathbb{R}^d \rightarrow R$ and $K(x, x')$ $\widetilde{K}(x, x') = f(x)K(x, x')f(x')$ $\cdot K(x, x') = K_1(x, x') + K_2(x, x')$ $\cdot K(x, x') = K_1(x, x') + K_2(x, x')$ radial basis kernel $K(x, x') = \exp\left(-\frac{1}{2}\ x - x'\ ^2\right)$ \Rightarrow correctly classifies every point \Rightarrow radial in $\phi(x)$ space graph classification look at number of bends in figure to determine dimension	L1 norm $\ w\ _1 = \sum_i w_i $ L2 norm $\ w\ _2 = \sqrt{\sum_i^n w_i^2}$ cosine similarity $\cos\left(x^{(i)}, x^{(j)}\right) = \frac{x^{(i)} \cdot x^{(j)}}{\ x^{(i)}\ \ x^{(j)}\ }$ euclidean distance $dist\left(x^{(i)}, x^{(j)}\right) = \ x^{(i)} - x^{(j)}\ $
we want large boundaries: $\vdots \mid \vdots$ where \vdots are $\theta \cdot x^{(i)} + \theta_0 = \{1, -1\}$ and \mid is $\theta \cdot x^{(i)} + \theta_0 = 0$ \cdot prevents overfitting \cdot the bigger $\ \theta\ $ the faster (= shorter distance) these reach -1 and 1 \cdot we can control the margins distance with $\ \theta\ $			
signed distance $\gamma_i = y^{(i)} \frac{\theta \cdot x^{(i)} + \theta_0}{\ \theta\ } = \frac{1}{\ \theta\ }$ or $-\frac{1}{\ \theta\ }$ to maximise $\frac{1}{\ \theta\ }$ we minimise $\frac{1}{2} \ \theta\ ^2$			
hinge loss $Loss_h(z) = Loss_h(y^{(i)}(\theta \cdot x^{(i)} + \theta_0)) = \max(0, 1 - y^{(i)}(\theta \cdot x^{(i)} + \theta_0))$ $= \begin{cases} 0 & \text{if } z \geq 1 \quad \text{correct side, outside margins} \\ 1 - z & \text{if } z < 1 \end{cases}$		NONLINEAR CLASSIFICATION not linearly separable \Rightarrow map x to higher dimension kernels i.e. $K(x, x') = \phi(x) \cdot \phi(X') = (x \cdot x') + (x \cdot x')^2$ where $\phi(x) = (x_1, x_2, x_1^2, \sqrt{2}x_1x_2, x_2^2)^T$ \Rightarrow dot product is cheap to compute	
hinge loss improves by moving θ towards example			
objective $J(\theta, \theta_0) = \frac{1}{n} \sum Loss_h\left(y^{(i)}\left(\theta \cdot x^{(i)} + \theta_0\right)\right) + \frac{\lambda}{2} \ \theta\ ^2$ where λ is regularisation loss and minimising that term increases margins \rightarrow by minimising J we balance maximising margins and minimising loss	COLLABORATIVE FILTERING assumes X has low rank $J(u, v) = \sum_{a, i \in D} \frac{(Y_{ai} - u_a v_i)^2}{2} + \frac{\lambda}{2} \sum u_a^2 + \frac{\lambda}{2} \sum v_i^2$ $X = \begin{bmatrix} \text{user1rating1} & \text{user1rating2} \\ \text{user2rating1} & \text{user2rating2} \end{bmatrix} = u \cdot v^T$ \Rightarrow u: user's rating tendency \Rightarrow v: info about movie	K-NEAREST NEIGHBOURS $\hat{Y}_{ai} = \frac{1}{K} \sum_{b \in \text{KNN}(a, i)} Y_{bi}$ or weigh by similarity $\hat{Y}_{ai} = \frac{\sum_{b \in \text{KNN}(a, i)} sim(a, b) Y_{bi}}{\sum_{b \in \text{KNN}} sim(a, b)}$ euclidian distance $\ x_a - x_b\ $ cosine similarity $\cos(\theta) = \frac{x_a \cdot x_b}{\ x_a\ \ x_b\ }$ limitations my combination of tastes maybe uncommon	
SOLVING QUADRATIC		ACTIVATION FUNCTIONS	
if linearly separable and don't allow any errors, can solve quadratic: $\Rightarrow J(\theta, \theta_0) = \frac{1}{2} \ \theta\ ^2$	alternating projections set v, optimise u, set u... continue until local convergence (exercise link)	ReLU $f(z) = \max(0, z)$ tanh $f(z) = \tanh(z)$ linear $f(z) = z$	

DEEP LEARNING

- intermediate layer is like feature representation: "learning ϕ "
- activation layer acts like linear classifier


backpropagation

compute partial derivatives using chain rule to do gradient descent

$$w_1 \leftarrow w_1 - \eta \left(\frac{\delta \text{Loss}}{\delta w_1} \right)$$

$$\frac{\delta \text{Loss}}{\delta w_1} = \frac{\delta \text{Loss}}{\delta f_1} \cdot \frac{\delta f_1}{\delta z_1} \cdot \frac{\delta z_1}{\delta w_1}$$


risks are vanishing or exploding gradients with long chains



RNN

- same params for each layer
- variable number of layers
- ⇒ encodes sentences

$$s_t = \tanh(W^{s,s} s_{t-1} + W^{s,x} x_t)$$



RNN have some gradient issues

gating network

in default RNN we overwrite s_t but sometimes it's useful to retain control over what's overwritten at each step

$$g_t = \text{sigmoid}(W^{g,s} s_{t-1} + W^{g,x} x_t)$$
$$s_t = (1 - g_t) \odot s_{t-1} + g_t \odot \tanh(W^{s,s} s_{t-1} + W^{s,x} x_t)$$

where \odot is element-wise multiplication

$$(1 - g_t) : \begin{pmatrix} 0 \\ 1 \end{pmatrix} \rightarrow \begin{pmatrix} 1 \\ 0 \end{pmatrix} \text{ i.e. } g_t \text{ is use to control what we retain from previous layer}$$

MARKOV MODELS

<beg> w_1 w_2 w_3 <end>
 w_0 w_L

$$P(w_1, ..., w_L) = P(w_1 | w_0) ... P(w_L | w_{L-1})$$

$$\text{with } P(w' | w) = \frac{\text{count}(w, w')}{\sum_{\tilde{w}} \text{count}(w, \tilde{w})}$$

RNNs for sequences

variable history

$$\phi(\text{word}) : X^W : s_t X : p$$
$$s_{t-1} : X$$

i.e. we're feeding the previous state as input to the next


CNN

convolution

- classify by small patches
- slide over input
- use the same weights for each patch to build feature map

pooling (subsampling)

- stores if there's an activation for feature in that area
- stride



cross-correlation/convolution

we don't actually need to flip $g(t)$ as we would with convolution, because we're learning the filter, the values are just flipped

$$1d \ f \odot g(t) = \sum f(T)g(T + t)$$

$$md \ f \odot g(x, y) = \sum_{T_1} \sum_{T_2} f(T_1, T_2)g(T_1 + x, T_2 + y)$$

K-MEANS

- 1) randomly assign $z^{(1)} ... z^{(2)}$
- 2) iterate
 - a) assign x to closest z
 - b) given $C_1 ... C_k$ find best rep. z in each C

$$\text{Cost} \left(z^{(1)} ... z^{(k)} \right) = \sum_{i=1}^n \min_{j=1...k} \left\| x^{(i)} - z^{(j)} \right\|^2$$

$$\text{Cost}(C_1 ... C_k) = \min_{z^{(1)} ... z^{(k)}} \sum_{j=1}^k \sum_{i \in C_j} \left\| x^{(i)} - z^{(j)} \right\|^2$$

problems with K-means

- points are not in set of X_s
- (euclidean distance)² only
- K-means is $O(n \cdot k \cdot d)$ for one update step (2)

find best z_j for C_j

$$\frac{\delta}{\delta z_j} \sum_{i \in C_j} \left\| x^{(i)} - z^{(j)} \right\|^2 = 0$$

$$z^{(j)} = \frac{\sum x^{(i)}}{|C_j|}$$

where $|C_j|$ is size of C_j

impact of init

- init matters
- we can end up with undesirable convergence

K-MEDIANS

Like K-means but with medians. Compute median in each single dimension in the Manhattan distance formulation (see [ex.](#))

K-MEDOIDS

- 1) pick init $z^{(i)} ... z^{(j)}$ from points X
- 2) iterate until no change in cost
 - a) for $i = 1 ... n$

$$C_j = \left\{ i : z^{(j)} \text{ is closest to } x^{(i)} \right\}$$

- b) pick the most central $x^{(i)}$ in C_j to be $z^{(j)}$

for $j = 1 ... k$

$$z_j = \left\{ x^{(i)} ... x^{(n)} : \sum_{i \in C_j} \text{dist} \left(x^{(i)}, z^{(j)} \right) \text{ is minimal} \right\}$$

· same idea as K-means but z_j is in x and can use any distance measure

· K-medoids is $O(n^2 \cdot k \cdot d)$ for update step (2)

Δ step 2 gets executed even if no change and cost, and order we cycle in matters and can update the medoid (see [ex.](#))

GENERATIVE MODELS

$p(w | \theta) = \theta w$: probability of generating word w given all possibilities`

probability of generating doc D

$$p(D | \theta) = \prod_{i=1}^n \theta_{w_i} = \prod_{w \in W} \theta_w^{\text{count}(w)}$$

MLE

$$\max P(D | \theta) \rightarrow \sum_{w \in W} \text{count}(w) \log(\theta_w)$$

$$\Rightarrow \tilde{\theta}_w = \frac{\text{count}(w)}{\sum_{w' \in W} \text{count}(w')}$$

prediction

essentially acts as a linear classifier through origin

$$\log \frac{p(D | \theta^+)}{p(D | \theta^-)} = \begin{cases} \geq 0 & + \\ < 0 & - \end{cases}$$

prior, posterior, likelihood

for when we have same prior knowledge of document classification

$$P(y = + | D) = \frac{P(D | \theta^+) P(y = +)}{P(D)}$$

⇒ linear classifier with offset influenced by prior

Δ N-1 params in multinomial (i.e. remember the binomial uses 1 param for two values)

GAUSSIAN GENERATIVE MODELS	MARKOV DECISION PROCESS	RL
<p>$X \mid \mu, \sigma^2 \sim N(\mu, \sigma^2)$</p> <p>want to find μ and σ^2 that give highest likelihood to training data.</p> <p>$\Rightarrow \frac{\delta l_n}{\delta \mu} = 0 \text{ and } \frac{\delta l_n}{\delta \sigma^2} = 0$</p> <p>$\widehat{\mu} = \frac{1}{n} \sum x^{(i)} \text{ and } \widehat{\sigma^2} = \frac{1}{nd} \sum \ x^{(i)} - \mu\ ^2$</p>	<p>· states $s \in S$</p> <p>· actions $a \in A$</p> <p>· action dependent on transition probabilities</p> <p>$T(s, a, s') = P(s' \mid a, s)$</p> <p>$\sum_{s' \in S} T(s, a, s') = 1$</p> <p>$R(s, a, s')$: reward for starting s, doing a, ending s'</p>	<p>typically won't have T and R ahead of time</p> <p>Q-value iteration</p> <p>want to calculate estimates for Q but don't have T and R. we'll use exponential running average to get estimates for $Q(s, a)$:</p> <p>$Q_{i+1}(s, a) = \alpha \cdot \text{sample} + (1 - \alpha)Q_i(s, a)$</p> <p>sample: $R(s, a, s') + \gamma \max_{a'} Q_i(s', a')$</p>
<p>MIXTURE MODELS</p> <p>mixture components K: $N(x, \mu_j, \sigma_j^2)$</p> <p>mixture weights $p_1 \dots p_k \sum p_j = 1$</p> <p>$\theta = p_1 \dots p_k, \mu_1 \dots \mu_k, \sigma_1^2 \dots \sigma_k^2$</p> <p>$p(x \mid \theta) = \sum_{j=1}^k p_j N(x, \mu_j, \sigma_j^2)$</p> <p>$p(S_n \mid \theta) = \prod_{i=1}^n \prod_{j=1}^k p_j N(x_i, \mu_j, \sigma_j^2)$</p>	<p>UTILITY FUNCTIONS</p> <p>finite horizon (not important)</p> <p>count reward for next n steps: $U(s_0 \dots s_n) = \sum^n R(s_i)$</p> <p>infinite horizon/discounted reward</p> <p>$U(s_0 \dots s_n) = \sum_{k=0}^\infty \gamma^k R(s_k) \leq \frac{R_{\max}}{1 - \gamma}$</p> <p>value function</p> <p>$V^*(s)$: expected reward if agent acts optimally starting at s</p> <p>$\pi: S \rightarrow A$: assigns action $\pi(s)$ to state s</p> <p>optimal policy π^* assigns action at every state that maximises expected utility</p>	<p>1) initialisation $Q(s, a) = 0 \forall a$</p> <p>2) iterate until convergence</p> <p>a) collect samples: $s, a, s', R(s, a, s')$</p> <p>b) $Q_{i+1}(s, a) = \alpha \left[R(s, a, s') + \gamma \max_a Q_i(s', a') \right] + (1 - \alpha)Q_i(s, a)$</p> <p>$\epsilon$ greedy</p> <p>balance exploration and exploitation</p> <p>· random action w.p. ϵ</p> <p>· best current action w.p. $(1 - \epsilon)$</p> <p>ϵ should decay over time</p>
<p>EM algorithm</p> <p>"random" init of $\theta: \mu^{(i)}, \sigma^{(i)}, p_i$</p> <p>init matters for speed and outcome</p> <p>· converges locally</p> <p>· good idea to use k-means for init</p> <p>expectation step</p> <p>$p(j \mid i) = \frac{p_j N(x_i, \mu_j, \sigma_j^2 I)}{p(x \mid \theta)}$</p> <p>$\Rightarrow$ probability that point i in cluster j</p> <p>$\Rightarrow p(x \mid \theta) = \sum_{j=1}^k p_j N(x_i, \mu_j, \sigma_j^2 I)$</p> <p>maximisation step</p> <p>re-estimate parameters, (nb: cluster j, point i)</p> <p>$\widehat{n}_j = \sum_{i=1}^n p(j \mid i)$</p> <p>$\widehat{p}_j = \frac{\widehat{n}_j}{n}$</p> <p>$\widehat{\mu}_j = \frac{1}{\widehat{n}_j} \sum_{i=1}^n p(j \mid i) x^{(i)}$</p> <p>$\widehat{\sigma}_j^2 = \frac{1}{\widehat{n}_j d} \sum_{i=1}^n p(j \mid i) \ x^{(i)} - \mu^{(j)}\ ^2$</p>	<p>$Q^*(s, a)$: expected reward starting at s, taking action a and acting optimally</p> <p>BELLMAN EQUATIONS</p> <p>$V^*(s) = \max_a Q^*(s, a) = Q^*(s, \pi^*(s))$</p> <p>$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')]$</p> <p>VALUE ITERATION ALGORITHM</p> <p>$V_k^*(s) \rightarrow_{k \rightarrow \infty} V(s)$</p> <p>· init: $V_0^*(s) = 0$</p> <p>· iterate until $V_k^*(s) \approx V_{k+1}^*(s) \forall s$</p> <p>$V_{k+1}^* = \max_a \left[\sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k^*(s')] \right]$</p> <p>· compute $Q^*(s, a)$</p> <p>then $\pi^*(s) = \arg \max_a Q^*(s, a)$</p> <p>Q-VALUE ITERATION</p> <p>same as above but with Q and the update is:</p> <p>$Q_{k+1}^* = \sum_{s'} T(s, a, s') \left[R(s, a, s') + \gamma \max_{a'} Q_k^*(s', a') \right]$</p> <p>$\Rightarrow$ we're summing over all possible states we may land in</p>	